

# Chapter 8

## I/O

## **I/O: Connecting to Outside World**

**So far, we've learned how to:**

- **compute with values in registers**
- **load data from memory to registers**
- **store data from registers to memory**

**But where does data in memory come from?**

**And how does data get out of the system so that humans can use it?**

# I/O: Connecting to the Outside World

## Types of I/O devices characterized by:

- **behavior:** input, output, storage
  - input: keyboard, motion detector, network interface
  - output: monitor, printer, network interface
  - storage: disk, CD-ROM
- **data rate:** how fast can data be transferred?
  - keyboard: 100 bytes/sec
  - disk: 30 MB/s
  - network: 1 Mb/s - 1 Gb/s

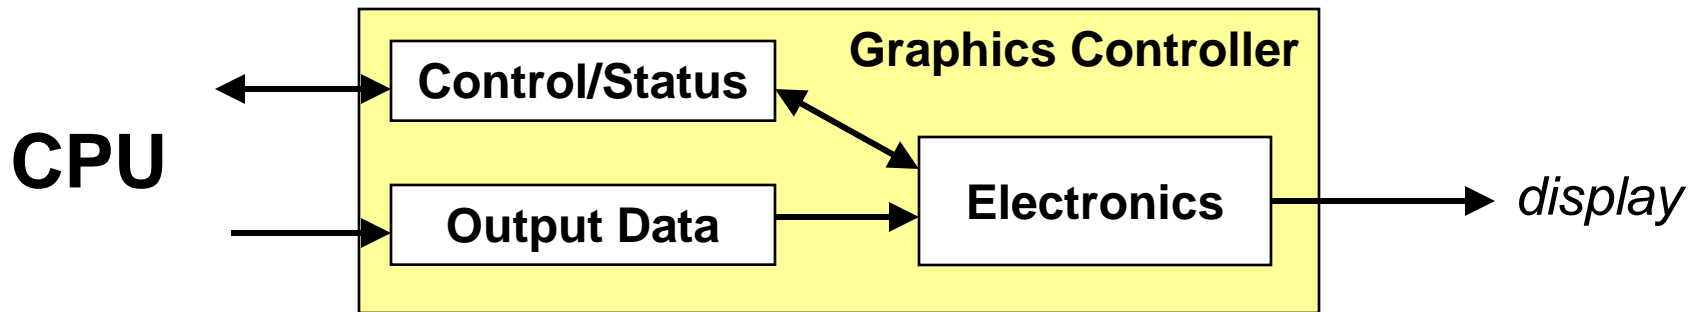
# I/O Controller

## Control/Status Registers

- CPU tells device what to do -- write to control register
- CPU checks whether task is done -- read status register

## Data Registers

- CPU transfers data to/from device



## Device electronics

- performs actual operation
  - pixels to screen, bits to/from disk, characters from keyboard

# Programming Interface

**How are device registers identified?**

- **Memory-mapped** vs. **special instructions**

**How is timing of transfer managed?**

- **Asynchronous** vs. **synchronous**

**Who controls transfer?**

- CPU (**polling**) vs. device (**interrupts**)

# Memory-Mapped vs. I/O Instructions


## Instructions

- designate opcode(s) for I/O
- register and operation encoded in instruction



## Memory-mapped

- assign a memory address to each device register
- use data movement instructions (LD/ST) for control and data transfer



## Transfer Timing

**I/O events generally happen much slower than CPU cycles.**

### Synchronous

- data supplied at a fixed, predictable rate
- CPU reads/writes every X cycles

### Asynchronous

- data rate less predictable
- CPU must synchronize with device, so that it doesn't miss data or write too quickly

# Transfer Control

Who determines when the next data transfer occurs?

## Polling

- CPU keeps checking status register until new data arrives OR device ready for next data
- “Are we there yet? Are we there yet? Are we there yet?”

## Interrupts

- Device sends a special signal to CPU when new data arrives OR device ready for next data
- CPU can be performing other tasks instead of polling device.
- “Wake me when we get there.”



## LC-3

### Memory-mapped I/O (Table A.3)

<i>Location</i>	<i>I/O Register</i>	<i>Function</i>
<b>xFE00</b>	Keyboard Status Reg (KBSR)	Bit [15] is one when keyboard has received a new character.
<b>xFE02</b>	Keyboard Data Reg (KBDR)	Bits [7:0] contain the last character typed on keyboard.
<b>xFE04</b>	Display Status Register (DSR)	Bit [15] is one when device ready to display another char on screen.
<b>xFE06</b>	Display Data Register (DDR)	Character written to bits [7:0] will be displayed on screen.

### Asynchronous devices

- synchronized through status registers

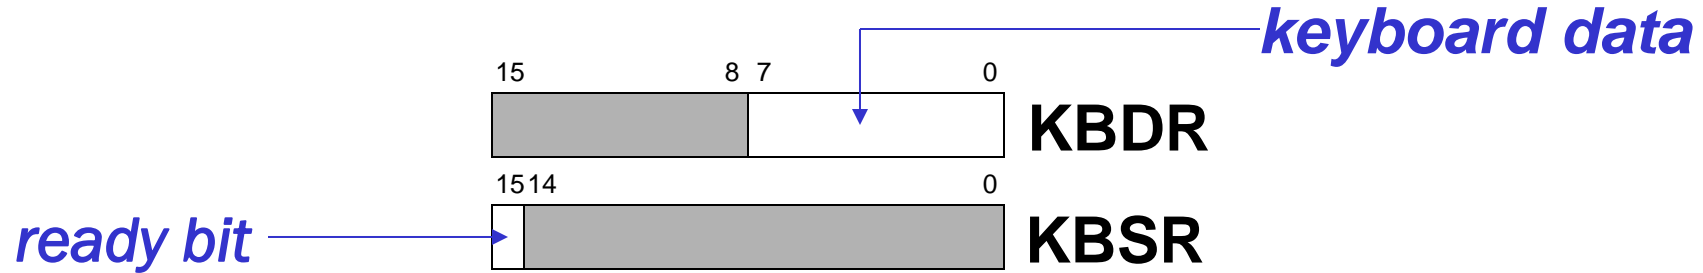
### Polling and Interrupts

- the details of interrupts will be discussed in Chapter 10

## Input from Keyboard

### When a character is typed:


- its ASCII code is placed in bits [7:0] of KBDR (bits [15:8] are always zero)
- the “ready bit” (KBSR[15]) is set to one
- keyboard is disabled -- any typed characters will be ignored



### When KBDR is read:

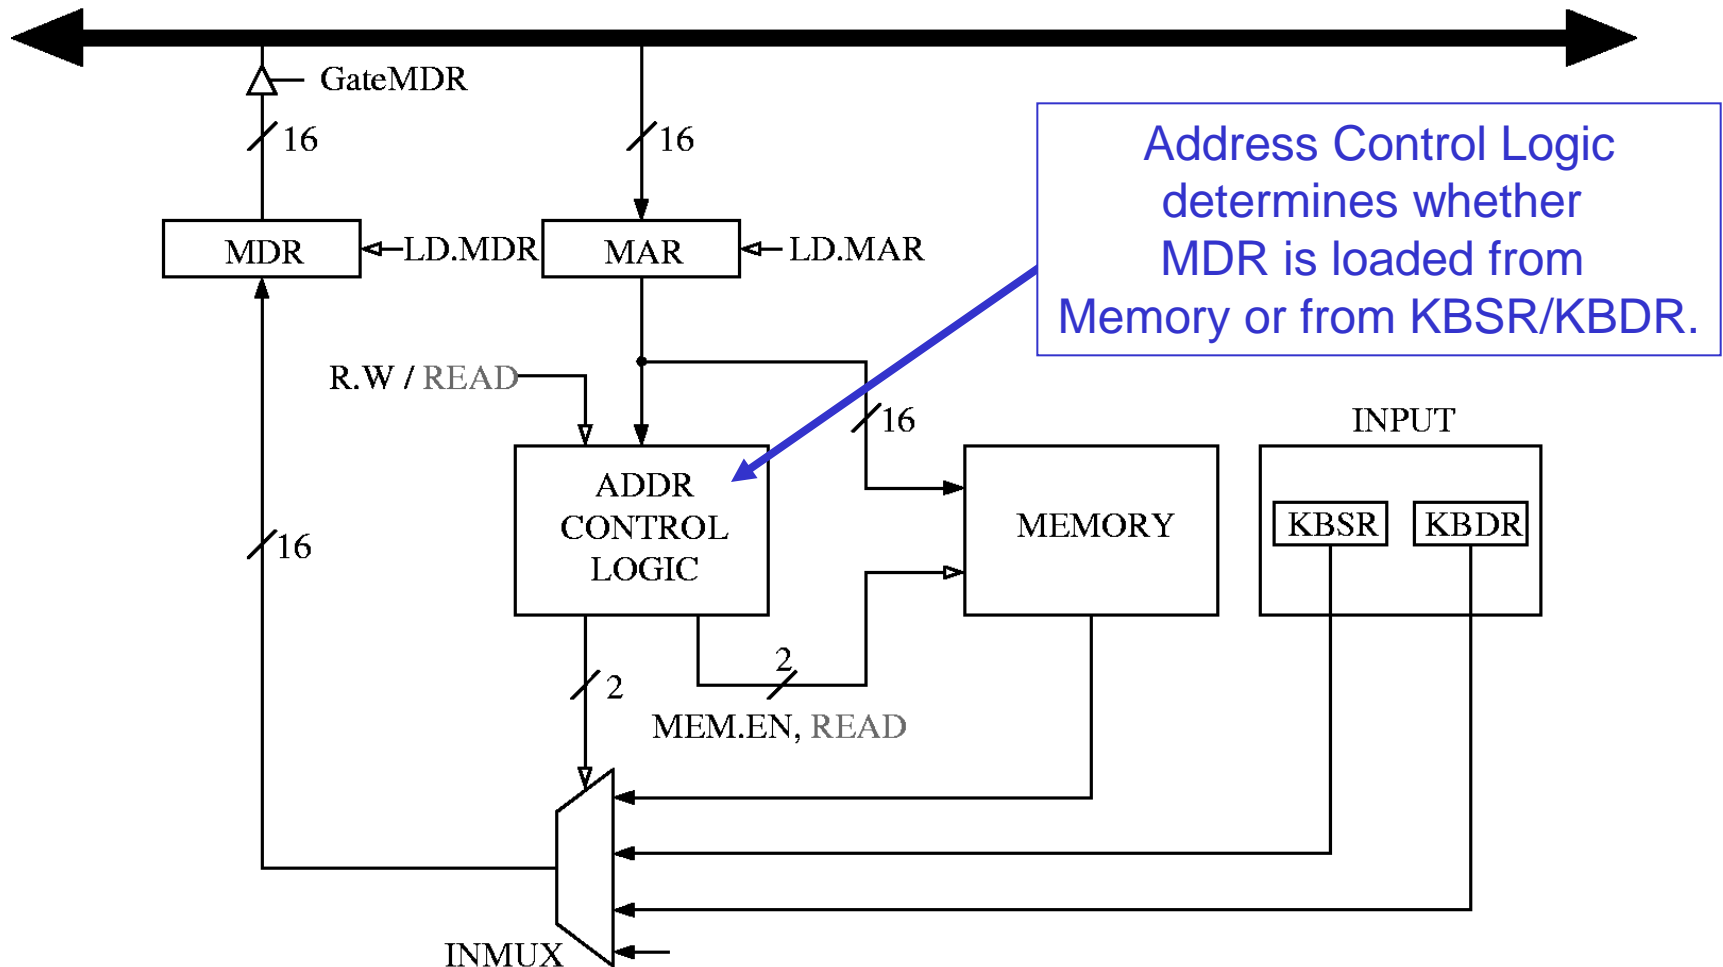
- KBSR[15] is set to zero
- keyboard is enabled

# Basic Input Routine



```
POLL      LDI    R0, KBSRPtr  
          BRzp   POLL  
          LDI    R0, KBDRPtr  
  
          . . .  
  
KBSRPtr   .FILL  xFE00  
KBDRPtr   .FILL  xFE02
```


# Simple Implementation: Memory-Mapped Input



## Output to Monitor

**When Monitor is ready to display another character:**

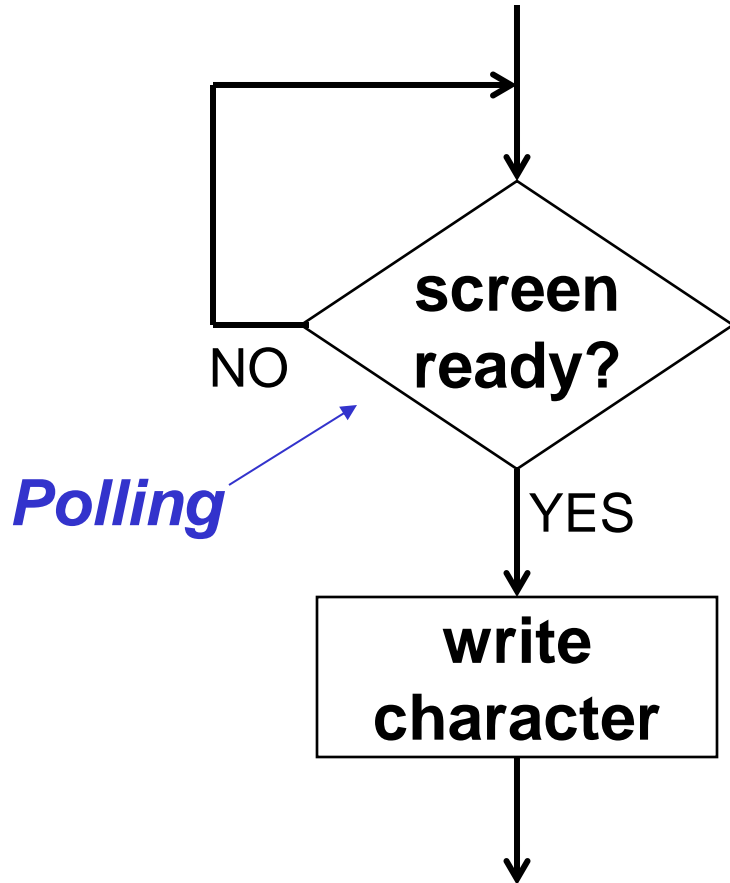
- the “ready bit” (DSR[15]) is set to one



**When data is written to Display Data Register:**

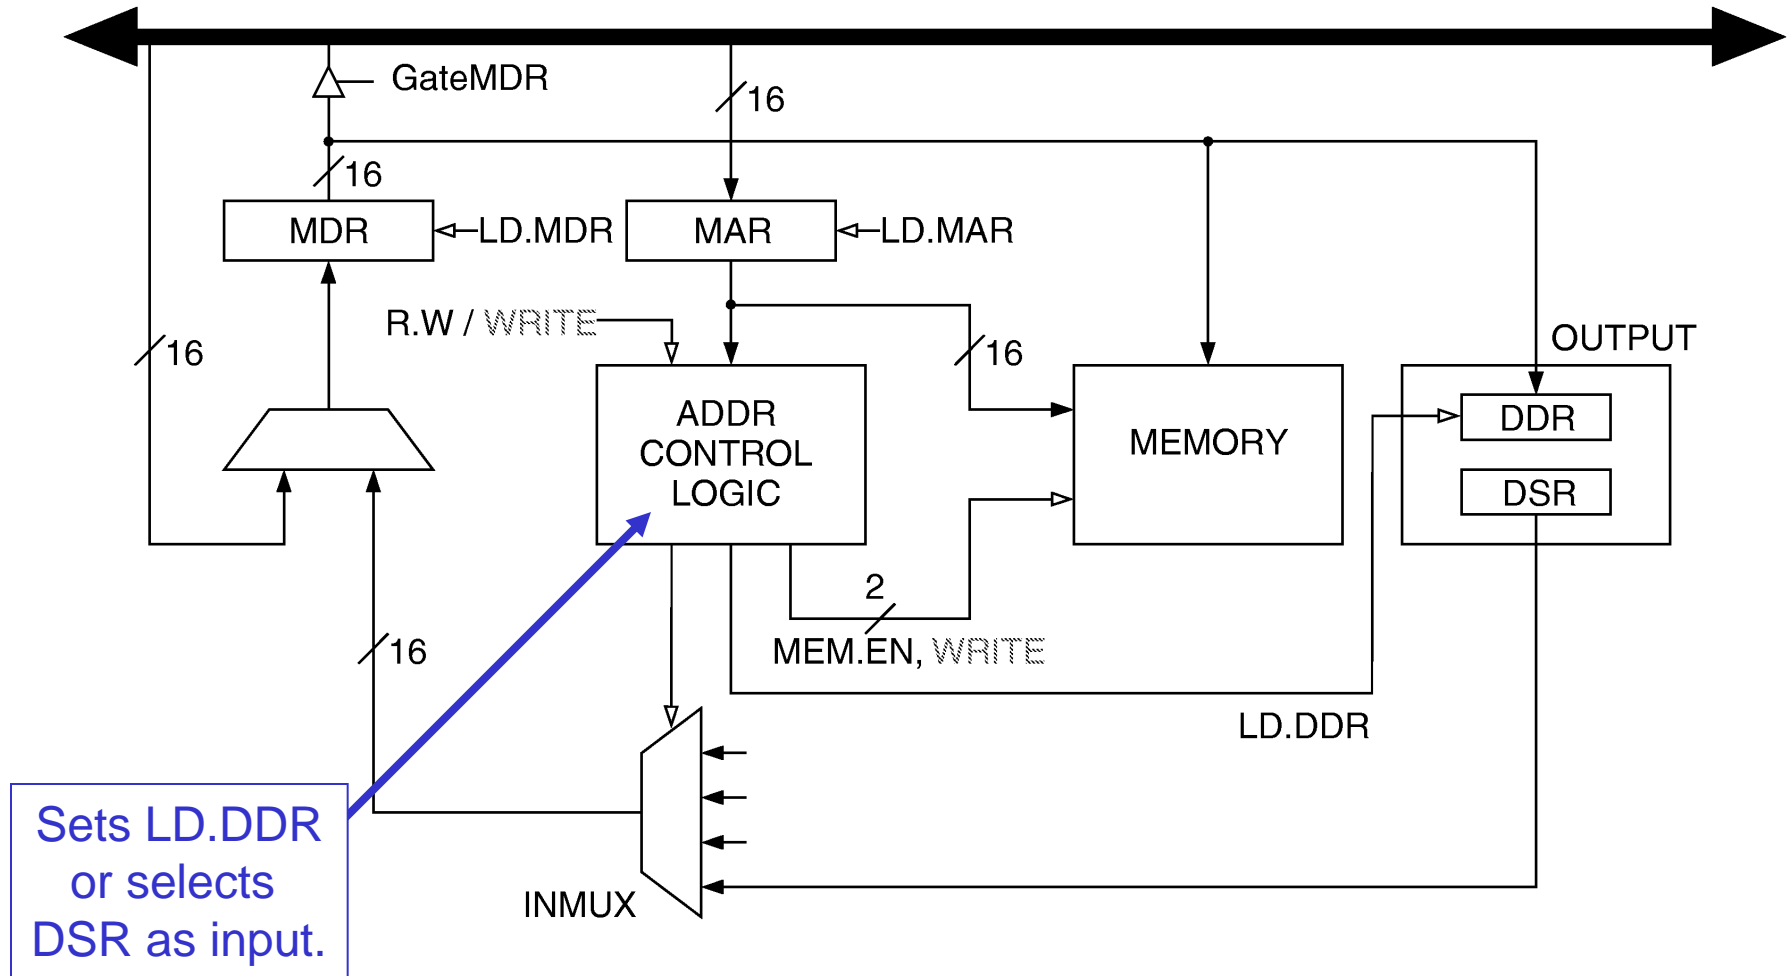
- DSR[15] is set to zero
- character in DDR[7:0] is displayed
- any other character data written to DDR is ignored (while DSR[15] is zero)

# Basic Output Routine



```
POLL    LDI    R1, DSRPtr  
        BRzp   POLL  
        STI    R0, DDRPtr  
  
        . . .  
  
DSRPtr  .FILL  xFE04  
DDRPtr  .FILL  xFE06
```

# Simple Implementation: Memory-Mapped Output



## Keyboard Echo Routine

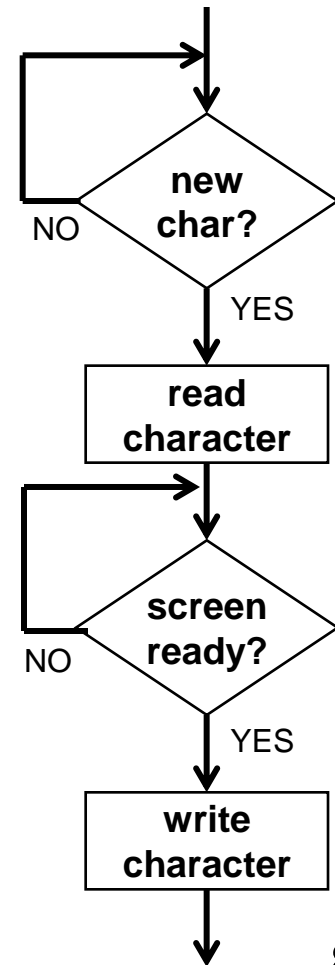
Usually, input character is also printed to screen.

- User gets feedback on character typed and knows its ok to type the next character.

```
POLL1    LDI    R0, KBSRPtr
          BRzp  POLL1
          LDI    R0, KBDRPtr
POLL2    LDI    R1, DSRPtr
          BRzp  POLL2
          STI    R0, DDRPtr

          . . .

KBSRPtr  .FILL  xFE00
KBDRPtr  .FILL  xFE02
DSRPtr   .FILL  xFE04
DDRPtr   .FILL  xFE06
```





## **Interrupt-Driven I/O**

**External device can:**

- (1) Force currently executing program to stop;**
- (2) Have the processor satisfy the device's needs; and**
- (3) Resume the stopped program as if nothing happened.**

**Why?**

- Polling consumes a lot of cycles, especially for rare events – these cycles can be used for more computation.**
- Example: Process previous input while collecting current input. (See Example 8.1 in text.)**

# **Interrupt-Driven I/O**

## **Generating the interrupt signal**

- **The I/O device must want to request service.**
- **The device must have the right to request service.**
- **This request must be more urgent than the processor's current task.**

## **Handling the interrupt signal**

- **We will wait until we understand stacks before getting to this. (Chapter 10)**


## Interrupt-Driven I/O

To implement an interrupt mechanism, we need:

- A way for the I/O device to **signal** the CPU that an interesting event has occurred.
- A way for the CPU to **test** whether the **interrupt signal is set** and whether its **priority is higher** than the current program.

### Generating Signal

- Software sets "interrupt enable" bit in device register.
- When ready bit is set and IE bit is set, interrupt is signaled.



# Generating the Interrupt Signal

## Using the Status Register

- The peripheral sets a Ready bit in SR[15] (as with polling)
- The CPU sets an Interrupt Enable bit in SR[14]
- These two bits are ANDed to set the Interrupt.  
(When ready bit is set and IE bit is set, interrupt is signaled.)

# Priority

**Every instruction is executed at a stated level of urgency.**

## LC-3: 8 priority levels (PL0-PL7)

- **Example:**
  - **Payroll program runs at PL0.**
  - **Nuclear power correction program runs at PL6.**
- **It's OK for PL6 device to interrupt PL0 program, but not the other way around.**

**Priority encoder** selects highest-priority device, compares to current processor priority level, and generates interrupt signal if appropriate.

# Priority

**Each task has an assigned priority level**

- **If a higher priority task requests access, a lower priority task will be suspended.**

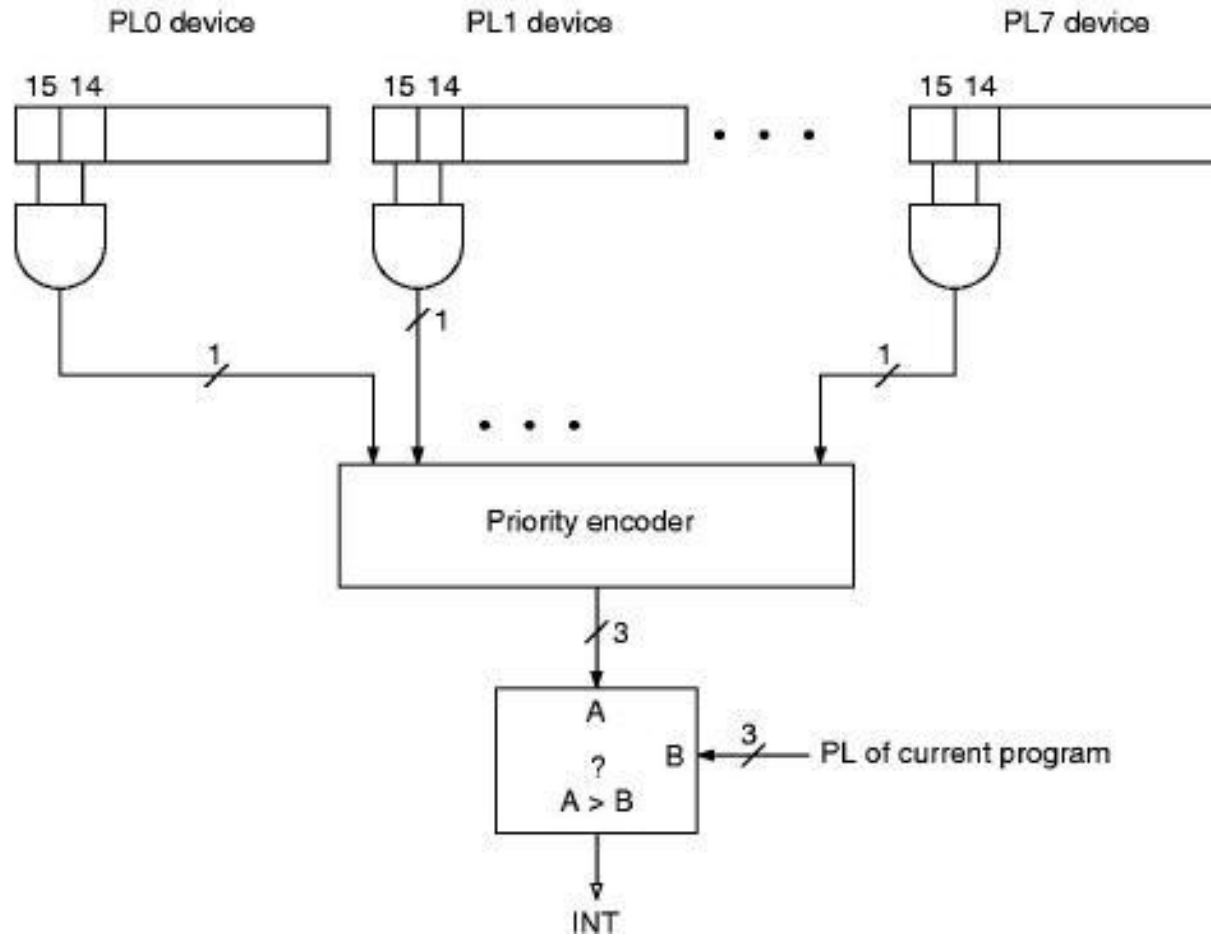
**Likewise, each I/O device request has an assigned priority**

- **The highest priority interrupt is passed on to the CPU only if it has higher priority than the currently executing task.**

**If an INT is present at the start of the instruction cycle, an extra step is inserted:**

- **The CPU saves its state information so that it can later return to the current task.**
- **The PC is loaded with the starting address of the Interrupt Service Routine**
- **The FETCH phase of the cycle continues as normal.**

# Generation of the INT Signal

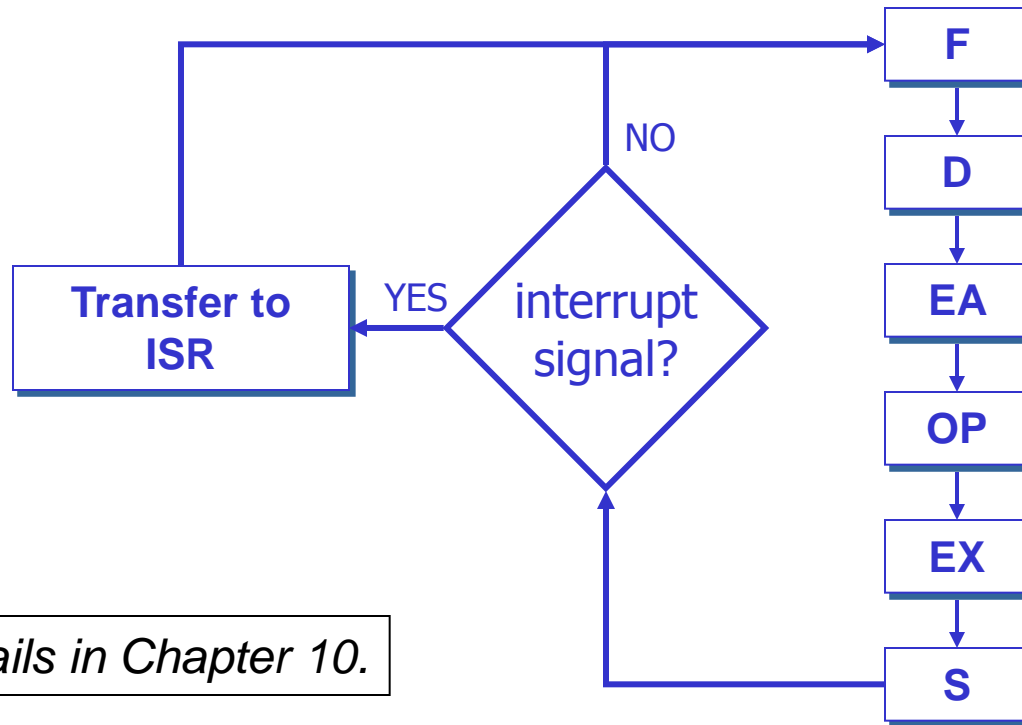


## Testing for Interrupt Signal

**CPU looks at signal between STORE and FETCH phases.**

**If not set, continues with next instruction.**

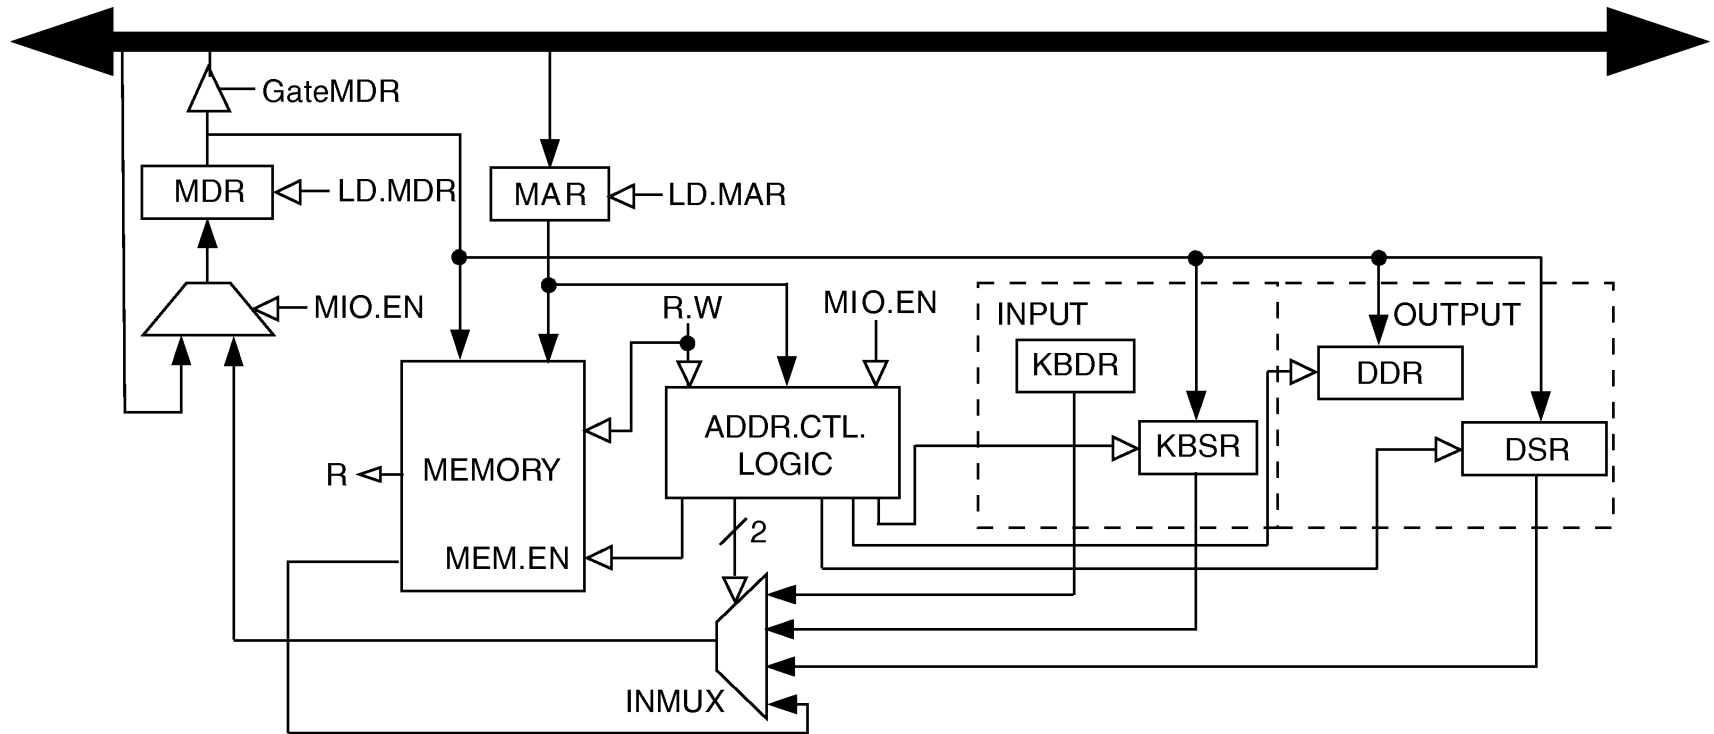
**If set, transfers control to interrupt service routine.**



*More details in Chapter 10.*



# Full Implementation of LC-3 Memory-Mapped I/O



Because of interrupt enable bits, status registers (KBSR/DSR) must be written, as well as read.

## **Review Questions**

**What is the danger of not testing the DSR before writing data to the screen?**

**What is the danger of not testing the KBSR before reading data from the keyboard?**

**What if the Monitor were a synchronous device, e.g., we know that it will be ready 1 microsecond after character is written.**

- **Can we avoid polling? How?**
- **What are advantages and disadvantages?**

## **Review Questions**

**Do you think polling is a good approach for other devices, such as a disk or a network interface?**

**What is the advantage of using LDI/STI for accessing device registers?**