

# Chapter 9

## TRAP Routines and Subroutines

## System Calls

Certain operations require **specialized knowledge** and **protection**:

- specific knowledge of I/O device registers and the sequence of operations needed to use them
- I/O resources shared among multiple users/programs; a mistake could affect lots of other users!

Not every programmer knows (or wants to know) this level of detail

Provide ***service routines*** or ***system calls*** (part of operating system) to safely and conveniently perform low-level, privileged operations

## System Call

1. User program invokes system call.
2. Operating system code performs operation.
3. Returns control to user program.

In LC-3, this is done through the ***TRAP mechanism***.

# LC-3 TRAP Mechanism

## ***1. A set of service routines.***

- **part of operating system -- routines start at arbitrary addresses**  
(convention is that system code is below x3000)
- **up to 256 routines**

## ***2. Table of starting addresses.***

- stored at **x0000** through **x00FF** in memory
- called **System Control Block** in some architectures


## ***3. TRAP instruction.***

- **used by program to transfer control to operating system**
- **8-bit trap vector names one of the 256 service routines**

## ***4. A linkage back to the user program.***

- **want execution to resume immediately after the TRAP instruction**

## TRAP Instruction



### Trap vector

- identifies which system call to invoke
- 8-bit index into table of service routine addresses
  - in LC-3, this table is stored in memory at **0x0000 – 0x00FF**
  - 8-bit trap vector is zero-extended into 16-bit memory address

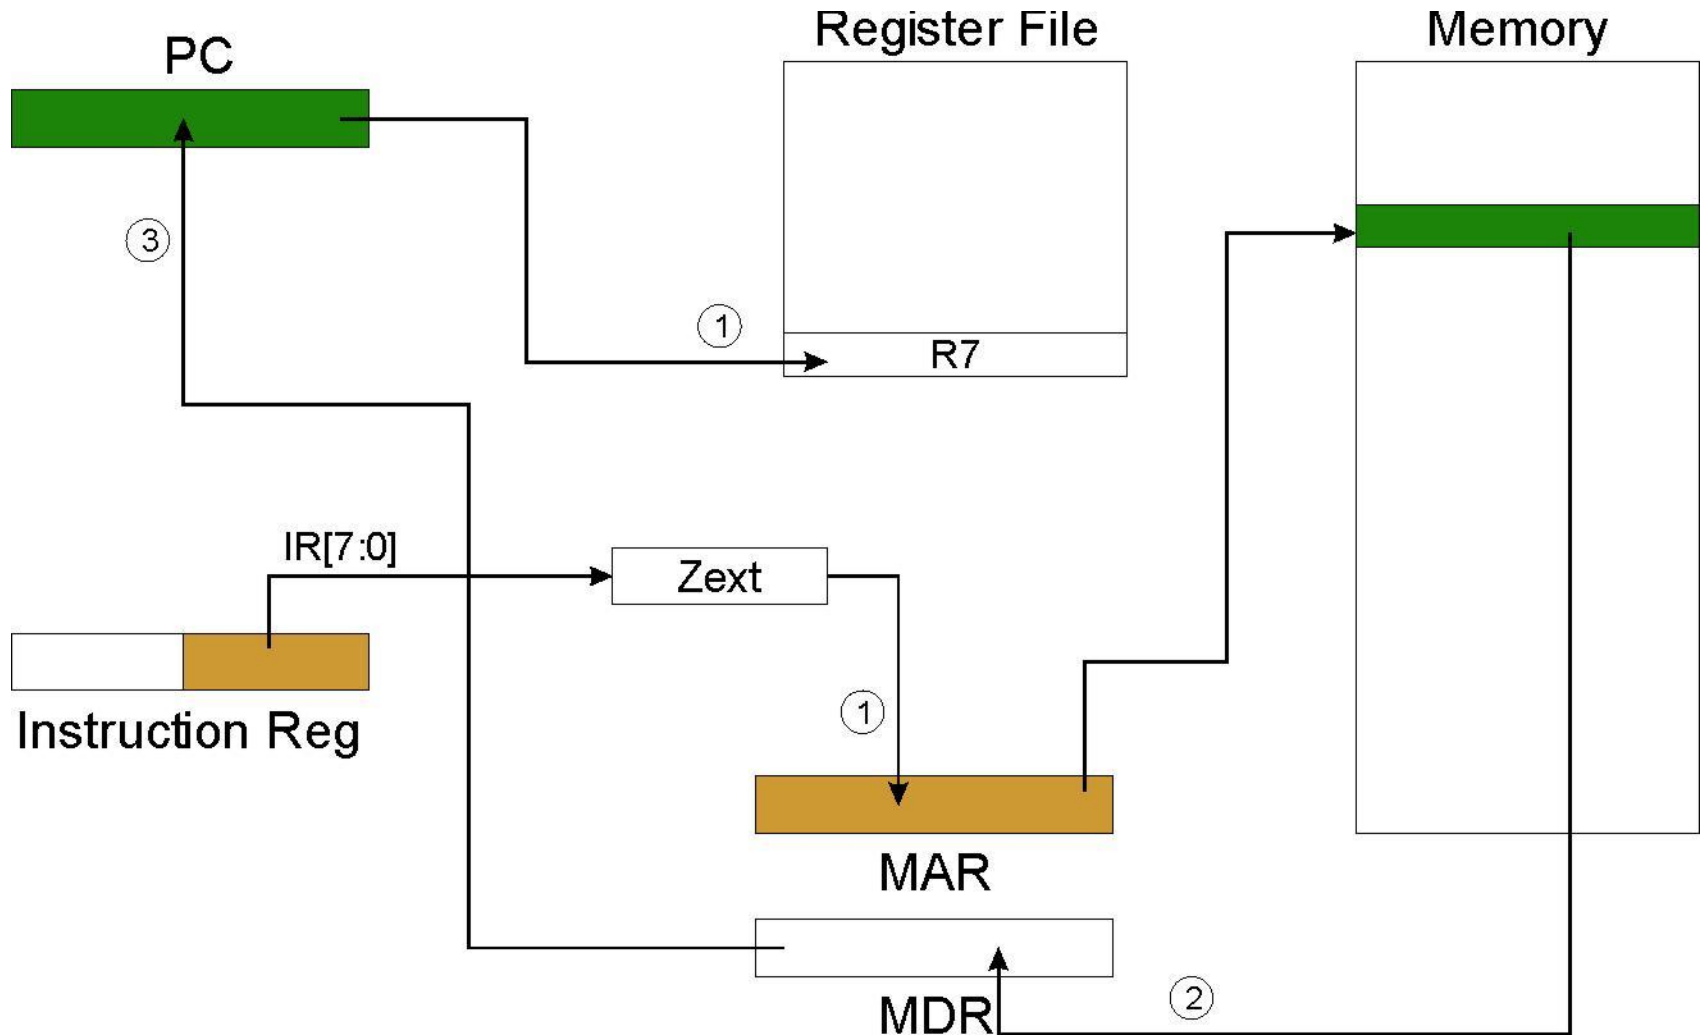
### Where to go

- lookup starting address from table; place in PC

### How to get back

- save address of next instruction (current PC) in R7

# TRAP



NOTE: PC has already been incremented during instruction fetch stage.

## **RET (JMP R7)**

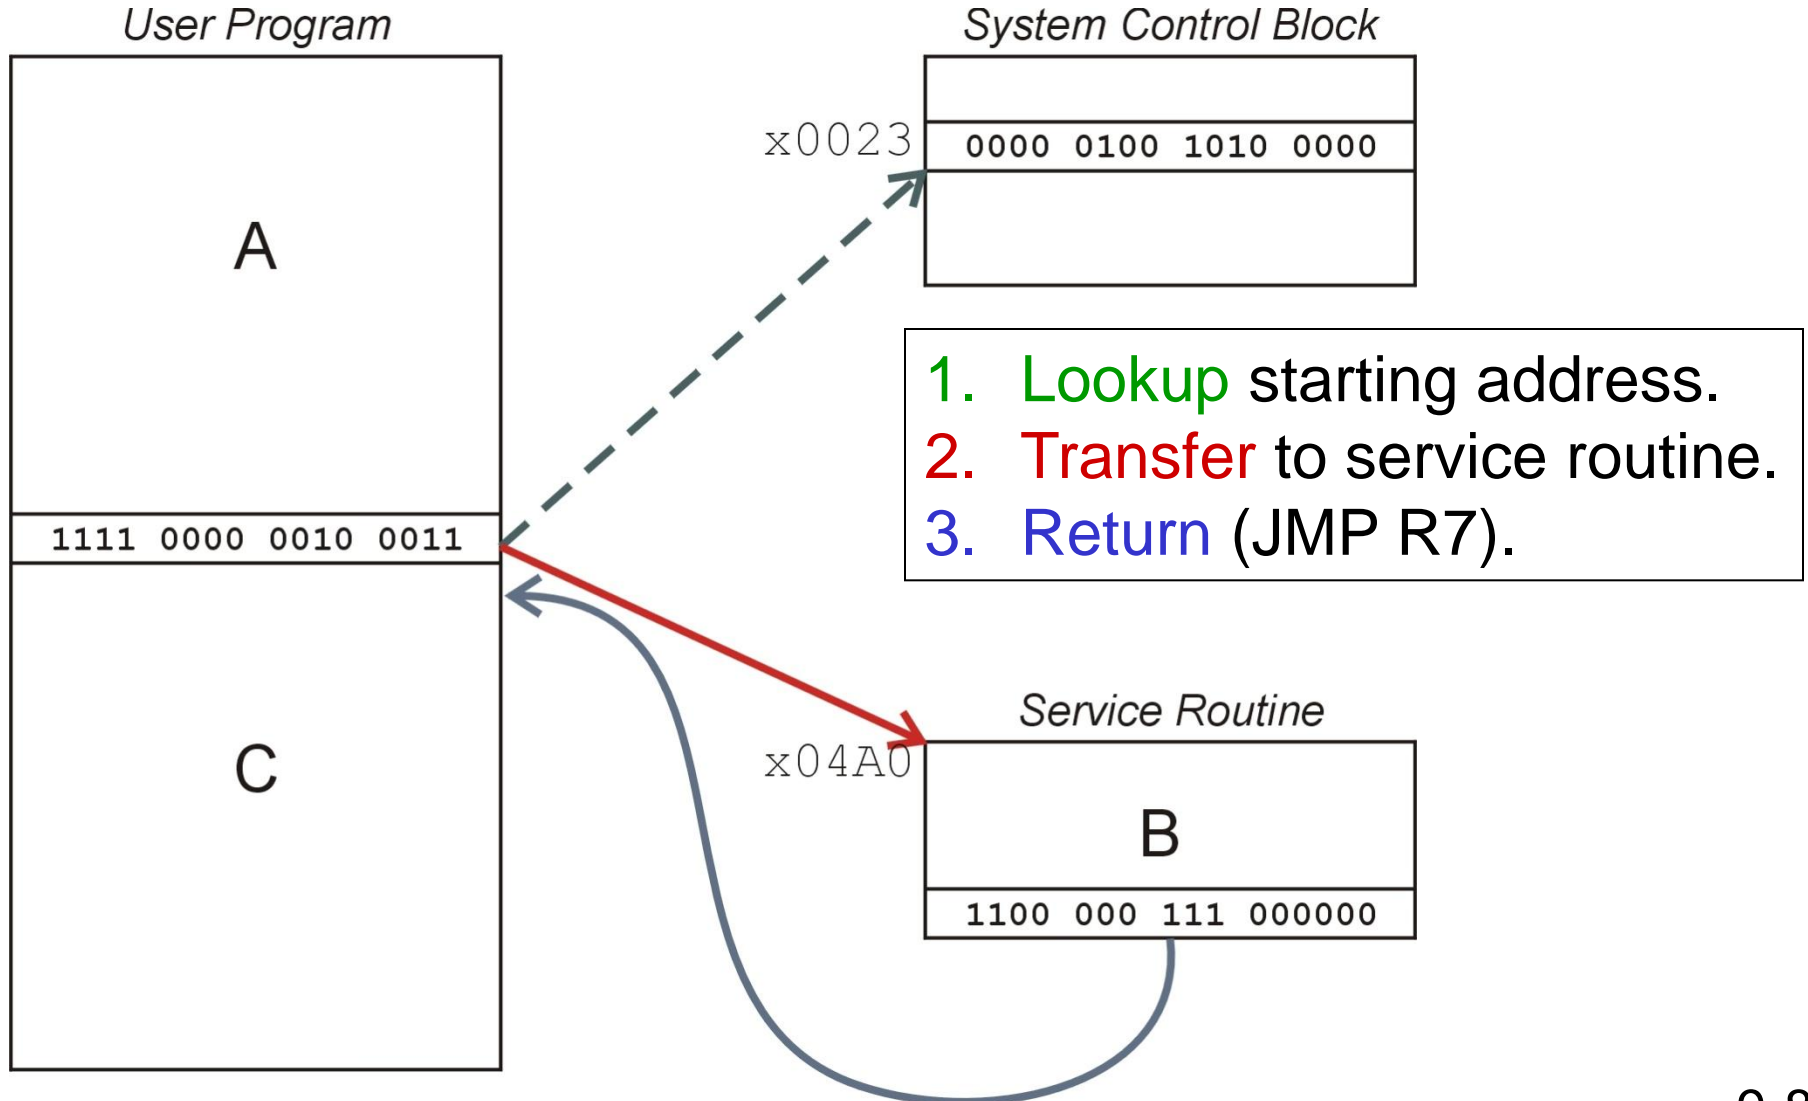
**How do we transfer control back to instruction following the TRAP?**

**We saved old PC in R7.**

- **JMP R7** gets us back to the user program at the right spot.
- **LC-3 assembly language** lets us use **RET** (return) in place of “JMP R7”.

**Must make sure that service routine does not change R7, or we won't know where to return.**

# TRAP Mechanism Operation





## Example: Using the TRAP Instruction

```
.ORIG x3000

LD    R2, TERM    ; Load negative ASCII '7'
LD    R3, ASCII    ; Load ASCII difference
AGAIN TRAP x23     ; input character
ADD   R1, R2, R0   ; Test for terminate
BRz   EXIT         ; Exit if done
ADD   R0, R0, R3   ; Change to lowercase
TRAP  x21          ; Output to monitor...
BRnzp AGAIN       ; ... again and again...

TERM  .FILL xFFC9  ; -'7'
ASCII .FILL x0020  ; lowercase bit
EXIT  TRAP x25     ; halt
      .END
```

## Example: Character Output Service Routine

```
        .ORIG x0430                ; syscall address
        ST      R7, SaveR7         ; save R7 & R1
        ST      R1, SaveR1
; ----- Write character
TryWrite  LDI     R1, DSR           ; get status
          BRzp    TryWrite         ; look for bit 15 on
WriteIt   STI     R0, DDR          ; write char
; ----- Return from TRAP
Return    LD      R1, SaveR1       ; restore R1 & R7
          LD      R7, SaveR7
          RET                     ; back to user

DSR       .FILL  xFE04
DDR       .FILL  xFE06
SaveR1    .FILL  0
SaveR7    .FILL  0
          .END
```

stored in table,  
location x21

# TRAP Routines and their Assembler Names

<i><b>vector</b></i>	<i><b>symbol</b></i>	<i><b>routine</b></i>
<b>x20</b>	<b>GETC</b>	read a single character (no echo)
<b>x21</b>	<b>OUT</b>	output a character to the monitor
<b>x22</b>	<b>PUTS</b>	write a string to the console
<b>x23</b>	<b>IN</b>	print prompt to console, read and echo character from keyboard
<b>x25</b>	<b>HALT</b>	halt the program

# **Saving and Restoring Registers**

**Must save the value of a register if:**

- **Its value will be destroyed by service routine, and**
- **We will need to use the value after that action.**

## **Who saves?**

- **caller of service routine?**
  - **knows what it needs later, but may not know what gets altered by called routine**
- **called service routine?**
  - **knows what it alters, but does not know what will be needed later by calling routine**

## Example

```

                                LEA    R3, Binary
                                LD      R6, ASCII    ; char->digit template
                                LD      R7, COUNT    ; initialize to 10
AGAIN                          TRAP    x23          ; Get char
                                ADD     R0, R0, R6    ; convert to number
                                STR     R0, R3, #0    ; store number
                                ADD     R3, R3, #1    ; incr pointer
                                ADD     R7, R7, #-1   ; decr counter
                                BRp     AGAIN        ; more?
                                BRnzp   NEXT
ASCII                          .FILL   xFFD0
COUNT                        .FILL   #10
Binary                        .BLKW   #10
```

What's wrong with this routine?  
What happens to R7?

# Saving and Restoring Registers

## Called routine -- *“callee-save”*

- Before start, save any registers that will be altered (unless altered value is desired by calling program!)
- Before return, restore those same registers

## Calling routine -- *“caller-save”*

- Save registers destroyed by own instructions or by called routines (if known), if values needed later
  - save R7 before TRAP
  - save R0 before TRAP x23 (input character)
- Or avoid using those registers altogether

*Values are saved by storing them in memory.*

## **Service Routines**

**Service routines provide three main functions:**

- 1. Shield programmers from system-specific details.**
- 2. Write frequently-used code just once.**
- 3. Protect system resources from malicious/clumsy programmers.**

# Subroutines

A **subroutine** is a program fragment that:

- lives in user space
- performs a well-defined task
- is invoked (called) by another user program
- returns control to the calling program when finished

Like a service routine, but not part of the OS

- not concerned with protecting hardware resources
- no special privilege required

Reasons for subroutines:

- reuse useful (and debugged!) code without having to keep typing it in
- divide task among multiple programmers
- use vendor-supplied *library* of useful routines



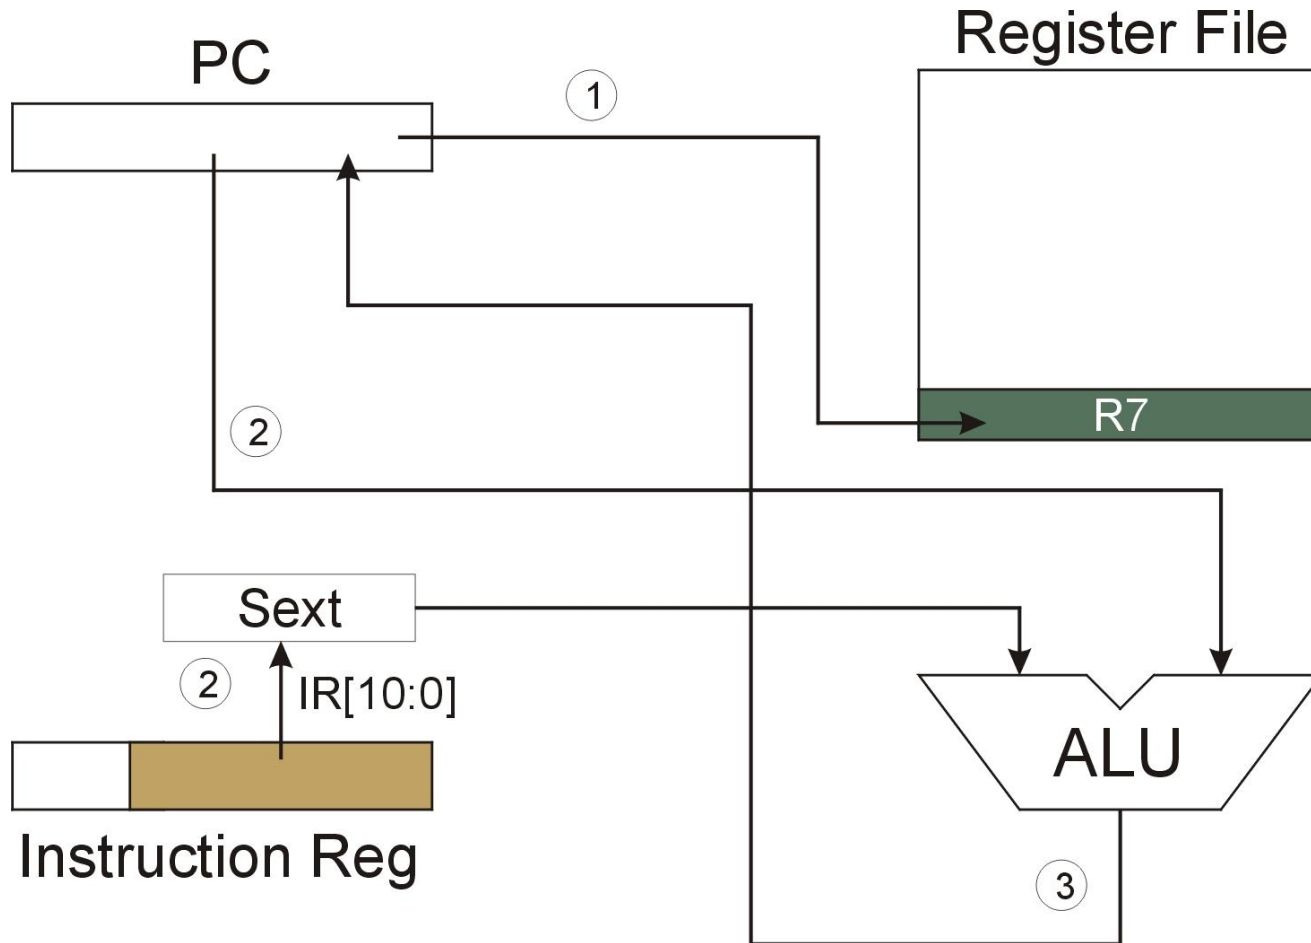
## JSR Instruction



**Jumps to a location (like a branch but unconditional), and saves current PC (addr of next instruction) in R7.**

- **saving the return address is called “linking”**
- **target address is PC-relative ( $PC + \text{Sext}(\text{IR}[10:0])$ )**
- **bit 11 specifies addressing mode**
  - **if =1, PC-relative: target address =  $PC + \text{Sext}(\text{IR}[10:0])$**
  - **if =0, register: target address = contents of register  $\text{IR}[8:6]$**

# JSR



NOTE: PC has already been incremented during instruction fetch stage.


## JSRR Instruction

	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
JSRR	0	1	0	0	0	0	0	Base			0	0	0	0	0	0

**Just like JSR, except Register addressing mode.**

- target address is Base Register
- bit 11 specifies addressing mode

# JSRR



NOTE: PC has already been incremented during instruction fetch stage.

## Returning from a Subroutine

**RET (JMP R7) gets us back to the calling routine.**

- just like TRAP

## Example: Negate the value in R0

```
2sComp    NOT    R0 , R0          ; flip bits
           ADD    R0 , R0 , #1    ; add one
           RET                                ; return to caller
```

***To call from a program (within 1024 instructions):***

```
; need to compute R4 = R1 - R3
           ADD    R0 , R3 , #0    ; copy R3 to R0
           JSR    2sComp          ; negate
           ADD    R4 , R1 , R0    ; add to R1
           . . .
```

***Note: Caller should save R0 if we'll need it later!***

# Passing Information to/from Subroutines

## Arguments

- A value **passed in** to a subroutine is called an argument.
- This is a value needed by the subroutine to do its job.
- Examples:
  - In 2sComp routine, R0 is the number to be negated
  - In OUT service routine, R0 is the character to be printed.
  - In PUTS routine, R0 is address of string to be printed.

## Return Values

- A value **passed out** of a subroutine is called a return value.
- This is the value that you called the subroutine to compute.
- Examples:
  - In 2sComp routine, negated value is returned in R0.
  - In GETC service routine, character read from the keyboard is returned in R0.

# Using Subroutines

**In order to use a subroutine, a programmer must know:**

- **its address** (or at least a label that will be bound to its address)
- **its function** (what does it do?)
  - **NOTE:** The programmer does not need to know how the subroutine works, but what changes are visible in the machine's state after the routine has run.
- **its arguments** (where to pass data in, if any)
- **its return values** (where to get computed data, if any)



## **Saving and Restoring Registers**

**Since subroutines are just like service routines, we also need to save and restore registers, if needed.**

**Generally use “callee-save” strategy, except for return values.**

- **Save anything that the subroutine will alter internally that shouldn't be visible when the subroutine returns.**
- **It's good practice to restore incoming arguments to their original values (unless overwritten by return value).**

**Remember: You MUST save R7 if you call any other subroutine or service routine (TRAP).**

- **Otherwise, you won't be able to return to caller.**

## Example

(1) Write a subroutine **FirstChar** to:

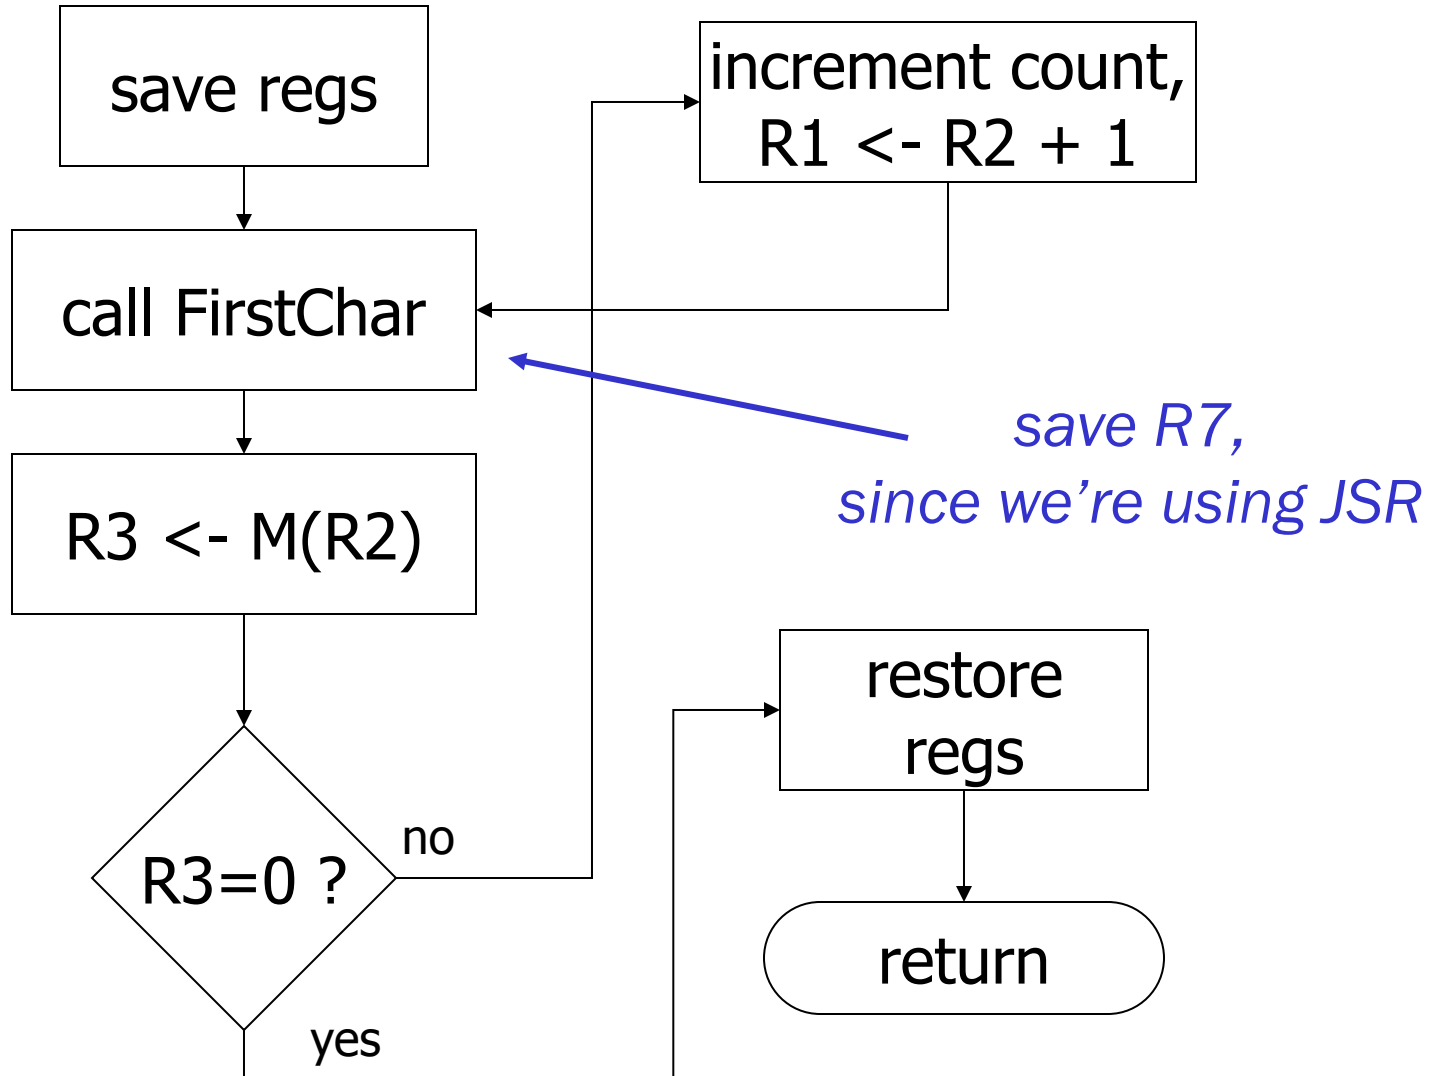
find the first occurrence  
of a particular **character** (in **R0**)  
in a **string** (pointed to by **R1**);  
return **pointer** to character or to end of string (NULL) in **R2**.

(2) Use **FirstChar** to write **CountChar**, which:

counts the number of occurrences  
of a particular **character** (in **R0**)  
in a **string** (pointed to by **R1**);  
return **count** in **R2**.

Can write the second subroutine first,  
without knowing the implementation of **FirstChar**!

## CountChar Algorithm (using FirstChar)



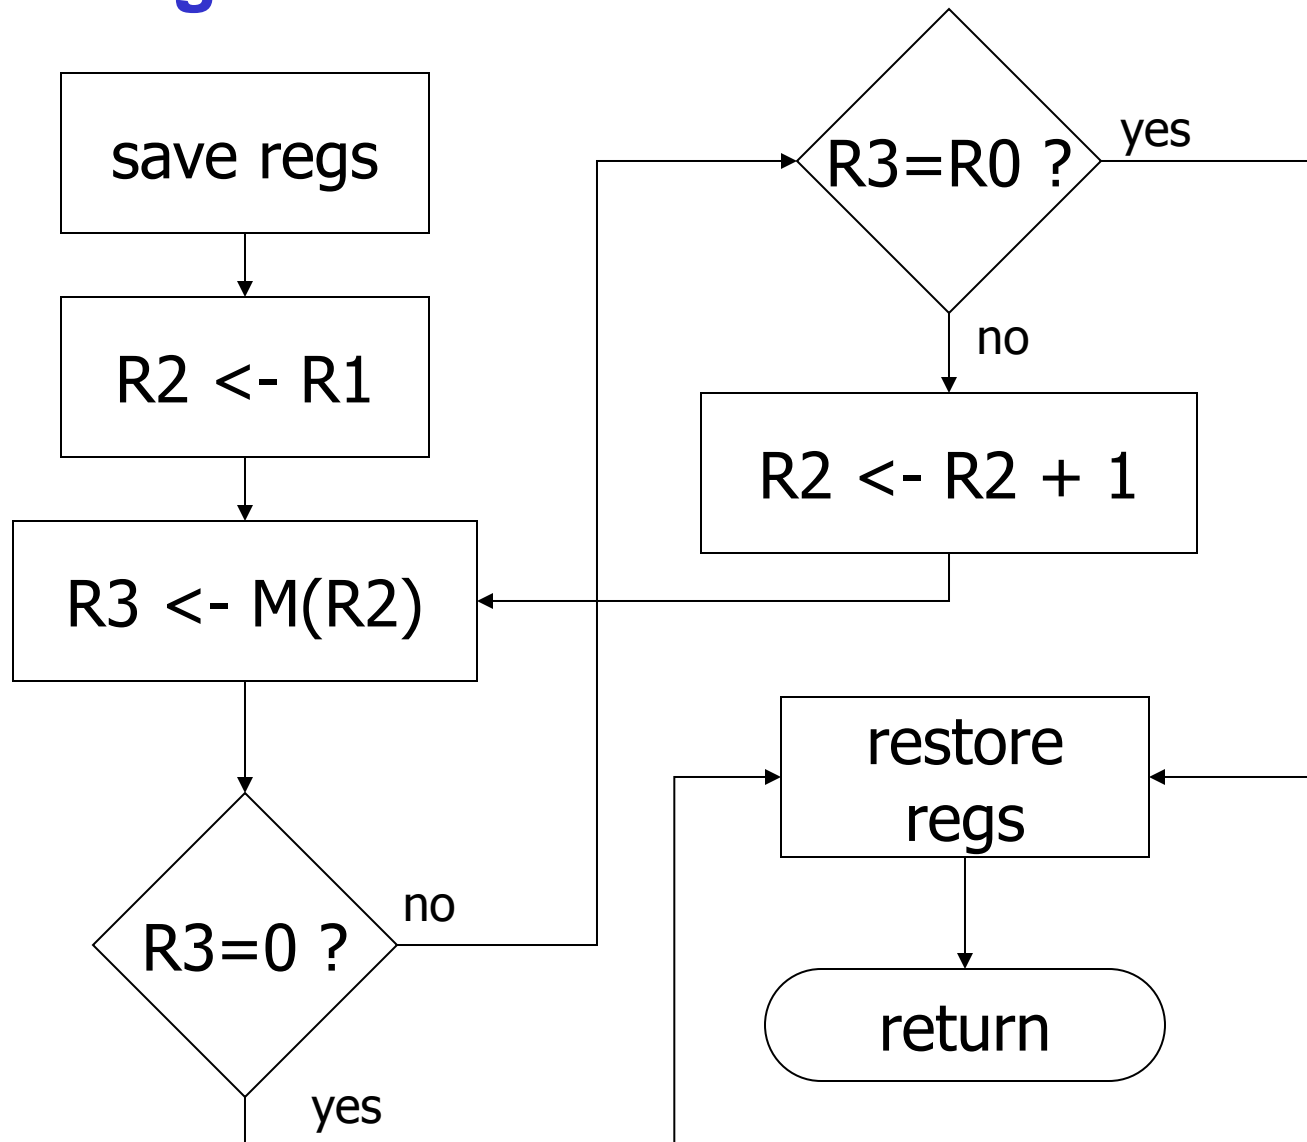
## CountChar Implementation

*; CountChar: subroutine to count occurrences of a char*

CountChar

```
        ST      R3, CCR3      ; save registers
        ST      R4, CCR4
        ST      R7, CCR7      ; JSR alters R7
        ST      R1, CCR1      ; save original string ptr
        AND     R4, R4, #0     ; initialize count to zero
CC1     JSR     FirstChar      ; find next occurrence (ptr in R2)
        LDR     R3, R2, #0     ; see if char or null
        BRz     CC2           ; if null, no more chars
        ADD     R4, R4, #1     ; increment count
        ADD     R1, R2, #1     ; point to next char in string
        BRnzp   CC1
CC2     ADD     R2, R4, #0      ; move return val (count) to R2
        LD      R3, CCR3      ; restore regs
        LD      R4, CCR4
        LD      R1, CCR1
        LD      R7, CCR7
        RET                               ; and return
```

# FirstChar Algorithm



## FirstChar Implementation

*; FirstChar: subroutine to find first occurrence of a char*

**FirstChar**

	<b>ST</b>	<b>R3 , FCR3</b>	<i>; save registers</i>
	<b>ST</b>	<b>R4 , FCR4</b>	<i>; save original char</i>
	<b>NOT</b>	<b>R4 , R0</b>	<i>; negate R0 for comparisons</i>
	<b>ADD</b>	<b>R4 , R4 , #1</b>	
	<b>ADD</b>	<b>R2 , R1 , #0</b>	<i>; initialize ptr to beginning of string</i>
<b>FC1</b>	<b>LDR</b>	<b>R3 , R2 , #0</b>	<i>; read character</i>
	<b>BRz</b>	<b>FC2</b>	<i>; if null, we're done</i>
	<b>ADD</b>	<b>R3 , R3 , R4</b>	<i>; see if matches input char</i>
	<b>BRz</b>	<b>FC2</b>	<i>; if yes, we're done</i>
	<b>ADD</b>	<b>R2 , R2 , #1</b>	<i>; increment pointer</i>
	<b>BRnzp</b>	<b>FC1</b>	
<b>FC2</b>	<b>LD</b>	<b>R3 , FCR3</b>	<i>; restore registers</i>
	<b>LD</b>	<b>R4 , FCR4</b>	<i>;</i>
	<b>RET</b>		<i>; and return</i>

## Library Routines

**Vendor may provide object files containing useful subroutines**

- don't want to provide source code -- intellectual property
- assembler/linker must support **EXTERNAL** symbols  
(or starting address of routine must be supplied to user)

```
. . .  
.EXTERNAL SQRT
```

```
. . .  
LD      R2, SQAddr      ; load SQRT addr  
JSRR    R2
```

```
. . .  
SQAddr  .FILL SQRT
```

**Using JSRR, because we don't know whether SQRT is within 1024 instructions.**